

Database Systems

Course Project Instruction

Ma Dongzhe / 马冬哲
152 1062 0224
mdzfirst@gmail.com

Policy

- 2-3 persons form a team.
- 60% of your final score.
- Submit at any time; test periodically.
- In December, 3-4 outstanding teams are invited to make a presentation.

The Task is

To Implement a DBMS Prototype.

What We Care

- Correctness
- Response Time
 - Storage
 - Access Method
 - Caching Strategy
 - Query Processing

What We Don't Care


- Transaction Processing
- Concurrency Control
- Crash Recovery

Grading Criteria

Accomplishment	At least one correct run	10
Overall Evaluation	Correctness & Design & Code Quality & Contrib.	10
Performance	$S_j = \frac{\sum((T_{i,best} / T_{i,j})^{0.2})}{Full} * (S_j / S_{best})$	30
Documentation	Content & Feature	10
Presentation	For some teams only	≤ 5

Example

	Workload 0	Workload 1
Team 0	5	100
Team 1	10	1000
Team 2	1	Fail



$$S_0 = (1 / 5)^{0.2} + (100 / 100)^{0.2} = 1.725$$

$$S_1 = (1 / 10)^{0.2} + (100 / 1000)^{0.2} = 1.262$$

$$S_2 = (1 / 1)^{0.2} + (100 / \text{INF})^{0.2} = 1$$

$$\text{Score}_0 = 30 * 1.725 / 1.725 = 30$$

$$\text{Score}_1 = 30 * 1.262 / 1.725 = 22$$

$$\text{Score}_2 = 30 * 1 / 1.725 = 17$$

$$(1 / 5)^{0.2} = 0.725$$

$$(1 / 10)^{0.2} = 0.631$$

$$(1 / 50)^{0.2} = 0.457$$

$$(1 / 100)^{0.2} = 0.398$$

$$(1 / 500)^{0.2} = 0.289$$

$$(1 / 1000)^{0.2} = 0.251$$

$$(1 / 5000)^{0.2} = 0.182$$

The Environment is

- Ubuntu 10.04 LTS, 32-bit
- g++ 4.4.3
- Intel(R) Xeon(R) 5130 @ 2.00GHz x2
- 3.0 GB RAM, 1.9 GB swap



Nonnegotiable

You Have to Implement

- **create()**

Create a new table.

- **train()**

Given some query information, train your system and choose the storage and access methods.

- **load()**

Load initial data in csv format. The initial data set might be too large to keep in the main memory entirely.

- **preprocess()**

Build the indexes and do other preprocessing.

See [course/include/client.h](#) for more details.

You Have to Implement

- `execute()`
Execute a query or insert statement.
- `next()`
Get the next row from the result set of the last query.
- `close()`
Close the sockets and kill other threads.

See `course/include/client.h` for more details.

You Have to Implement

- **execute()**

Execute a query or insert statement.

- **next()**

Get the next row from the result set of the last query.


- **close()**

Close the sockets and kill other threads.

**WARNING: Run time of execute()
and next() will be measured.**

See [course/include/client.h](#) for more details.

Test Procedure



Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```

A condition could be

```
column = constant  
column < constant (For integers only)  
column > constant (For integers only)  
column0 = column1 (Join condition)
```

Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```

A condition could be

No prefix

column = constant

column < constant (For integers only)

column > constant (For integers only)

column0 = column1 (Join condition)

Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```

A condition could be

column = constant

column < constant (For integers only)

column > constant (For integers only)

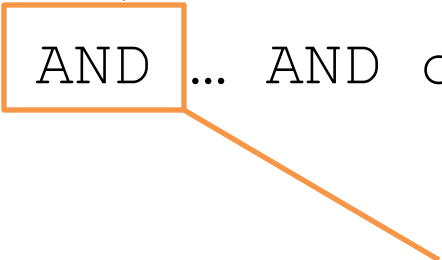
column0 = column1 (Join condition)



Same type

Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```



A condition could be

The only operator

column = constant

column < constant (For integers only)

column > constant (For integers only)

column0 = column1 (Join condition)

Query Statement

```
SELECT column0, column1, ...
```

```
FROM table0, table1, ...
```

```
WHERE condition0 AND .. AND conditionN;
```

A condition could be

If the FROM-clause contains only one table, there might be no WHERE-clause.

`column = constant`

`column < constant` (For integers only)

`column > constant` (For integers only)

`column0 = column1` (Join condition)

Insert Statement

```
INSERT INTO table  
VALUES (value_list0), ..., (value_listN);
```

All value lists are in csv format.

```
constant0,constant1,...,constantN
```

Insert Statement

```
INSERT INTO table  
VALUES (value_list0), ..., (value_listN);
```



All value lists are in csv format.

No column list

```
constant0, constant1, ..., constantN
```

Insert Statement

```
INSERT INTO table  
VALUES (value_list0), ..., (value_listN);
```

All value lists are in csv format. **Number of rows is important for the train() routine.**

```
constant0, constant1, ..., constantN
```

Data Types

- **INTEGER**

32-bit unsigned integer, 'int' is OK.

- **VARCHAR(d)**

*Consist of _, a-z, A-Z, or 0-9. Enclosed by single quotes.
At most d characters (excluding the quotes).*

NOTE:

All identifiers (table names and column names) are string constants not starting with 0-9.

Columns in different tables have distinct names.

String constants don't contain space, quote, or comma.

Primary Keys

- Primary keys will be assigned to all relations.
- The primary keys will be unique. There is no need to check this constraint.
- The primary keys will be given in ascending order.
- You can just ignore them.

Join Operations

Let nodes represent tables and edges represent join conditions, then each query can be transformed into a graph. This graph should be a tree which

- is connected;
- contains no self-cycles;
- contains no duplicate edges;
- contains no cycles (at least 3 nodes).

Workloads

Patience is a virtue.

For TPC-C and TPC-H, visit www.tpc.org.

Directory Structure

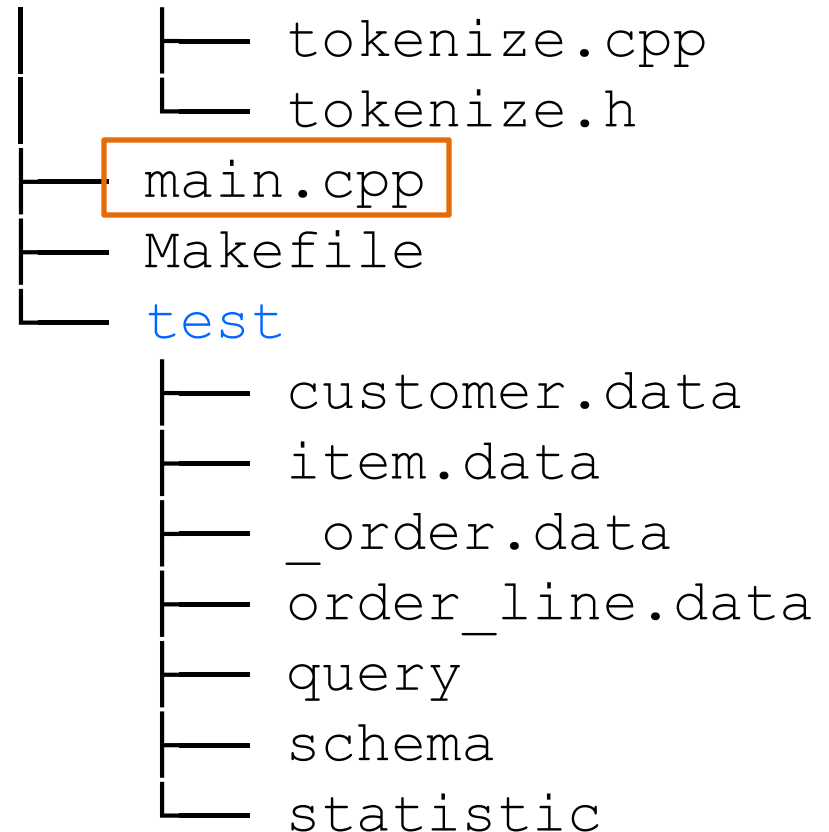
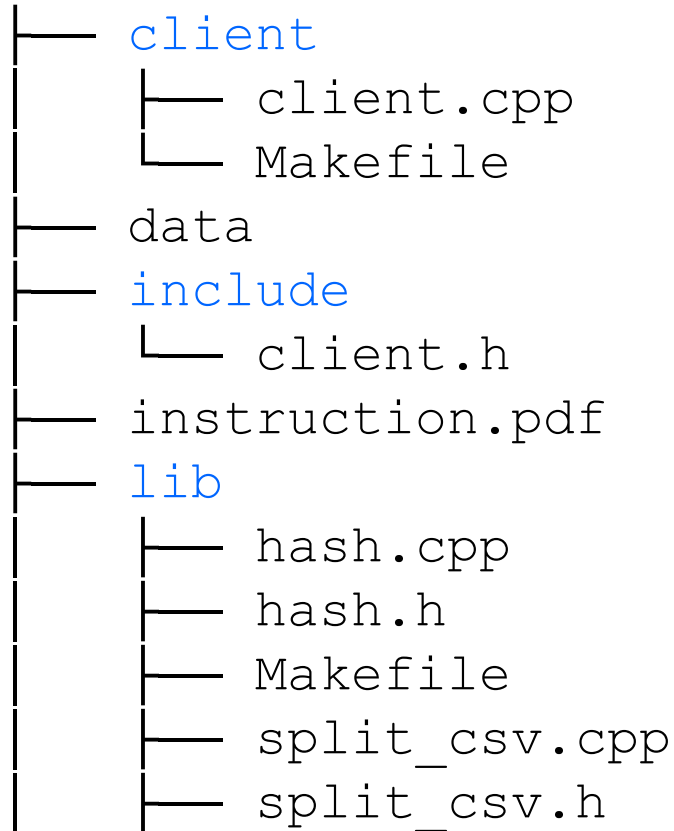
course

- client
 - client.cpp
 - Makefile
- data
- include
 - client.h
- instruction.pdf
- lib
 - hash.cpp
 - hash.h
 - Makefile
 - split_csv.cpp
 - split_csv.h

- tokenize.cpp
- tokenize.h
- main.cpp
- Makefile
- test
 - customer.data
 - item.data
 - _order.data
 - order_line.data
 - query
 - schema
 - statistic

Directory Structure

course



Test procedure. This file will be modified.

Directory Structure

course

- client
 - client.cpp
 - Makefile
- data
- include
 - client.h
- instruction.pdf
- lib
 - hash.cpp
 - hash.h
 - Makefile
 - split_csv.cpp
 - split_csv.h

- tokenize.cpp
 - tokenize.h
- main.cpp
- Makefile
- test
 - customer.data
 - item.data
 - _order.data
 - order_line.data
 - query
 - schema
 - statistic

APIs to implement.

Directory Structure

course

client

client.cpp

Makefile

data

include

client.h

instruction.pdf

lib

hash.cpp

hash.h

Makefile

split_csv.cpp

split_csv.h

tokenize.cpp

tokenize.h

main.cpp

Makefile

test

customer.data

item.data

_order.data

order_line.data

query

schema

statistic

Keep all your source codes in this directory.

Directory Structure

course

- client
 - client.cpp
 - Makefile
- data
- include
 - client.h
- instruction.pdf
- lib
 - hash.cpp
 - hash.h
 - Makefile
 - split_csv.cpp
 - split_csv.h

- tokenize.cpp
 - tokenize.h
- main.cpp
- Makefile
- test
 - customer.data
 - item.data
 - _order.data
 - order_line.data
 - query
 - schema
 - statistic

Make sure your Makefile is correct.

Directory Structure

course

- client
 - client.cpp
 - Makefile
- data
- include
 - client.h
- instruction.pdf
- lib
 - hash.cpp
 - hash.h
 - Makefile
 - split_csv.cpp
 - split_csv.h

- tokenize.cpp
- tokenize.h
- main.cpp
- Makefile
- test
 - customer.data
 - item.data
 - _order.data
 - order_line.data
 - query
 - schema
 - statistic

Keep all your data in this directory.

About Third-party Library

- You are free to use any third-party library or code about storage, index, multi-thread, network, etc.

e.g. Boost, Berkeley DB, open-source disk-based B-tree / hash table implementation, etc

- You are forbidden to use any system that is capable to process a SQL query.

e.g. MySQL, PostgreSQL, etc

- Ask for confirmation if you are not sure.

Document / Presentation

- System Architecture
- Storage Model and the Selection Strategy
- Index Structure and the Selection Strategy
- Caching Strategy
- Query Processing Strategy
 - Heuristic Rules
 - Cost Model
- Other features of your system
- References
- Personal Contribution Rate (For document)

Submission

- Send the compressed client directory to *mdzfirst@gmail.com* when you make a remarkable improvement.
- First come, first service.
- Only the last submission counts.
- Results will be made public periodically.

Hints

- Read some research papers
- Discuss with others
- Start ASAP

Warnings

- Never do irrelevant operations
- Never replicate other team's work



Any questions?

create()

Keep the schema safe.

train()

- Find affinitive tables.
- Find affinitive attributes.
- Choose access methods.
- Read-intensive or update-intensive?

Logical View

rid	name	birth	country
0	Stalin	1879	Soviet Union
1	Roosevelt	1882	United States
2	Churchill	1874	United Kingdom

[CODD1970] A Relational Model of Data for Large Shared Data Banks.

Horizontal Partitioning

1	Roosevelt	1882	United States
↓	↓	↓	↓
0	Stalin	1879	Soviet Union
2	Churchill	1874	United Kingdom

[CER1982] Horizontal Data Partitioning in Database Design.

Vertical Partitioning

0	Stalin	1879
1	Roosevelt	1882
2	Churchill	1874

0	Soviet Union
1	United States
2	United Kingdom

[NAVATHE1984] Vertical Partitioning Algorithms for Database Design.

DSM (MonetDB)

0	Stalin
1	Roosevelt
2	Churchill

0	1879
1	1882
2	1874

0	Soviet Union
1	United States
2	United Kingdom


Churchill	2
Roosevelt	1
Stalin	0

1874	2
1879	0
1882	1

Soviet Union	0
United Kingdom	2
United States	1

[COPELAND1985] A Decomposition Storage Model.
[KHO1987] A Query Processing Strategy for the Decomposed Storage Model.
[B] Monet: A Next-Generation DBMS Kernel for Query-Intensive Applications.
<http://www.monetdb.org/Home>

C-Store (Vertica)



[A] Query Execution in Column-Oriented Database Systems.
<http://db.csail.mit.edu/projects/cstore/>
<http://www.vertica.com/>

load()

Keep the data safe.

preprocess()

- Make some useful statistics.
- Build some indexes.
- Start some threads.

Statistics

- $\text{Size}(R)$, $\text{Cnt}(R)$, $\text{Card}(A)$, $\text{Min}(A)$, $\text{Max}(A)$
- $\text{SF}(A = \text{value}) = 1 / \text{Card}(A)$
- $\text{SF}(A > \text{value}) = (\text{Max}(A) - \text{value}) / \text{Range}(A)$
- $\text{SF}(A < \text{value}) = (\text{value} - \text{Min}(A)) / \text{Range}(A)$
- $\text{SF}(A_0 \wedge A_1) = \text{SF}(A_0) * \text{SF}(A_1)$

execute() and next()

- Do all the jobs in execute().
- Do all the jobs in next().
- Do all the jobs in independent threads.

Query Processing


- Google 'query processing'
- Search on ACM Digital Library (*dl.acm.org*)
- [GRAEFE1990] Encapsulation of Parallelism in the Volcano Query Processing System
- [GRAEFE1994] Volcano – An Extensible and Parallel Query Evaluation System

Join Operation

- Nested Loop Join
- Index-based Nested Loop Join
- Sort-Merge Join
- Hash Join (Pruning)

Past Contest

- 2009 Main Memory Transactional Index
db.csail.mit.edu/sigmod09contest/
- 2010 Distributed Query Engine
dbweb.enst.fr/events/sigmod10contest/
- 2011 A Durable Main-Memory Index
Using Flash
db.csail.mit.edu/sigmod11contest/



大学计算机教育国外著名教材、教参系列 (影印版)

Principles of Distributed Database Systems


PEARSON
Prentice
Hall

M. Tamer Özsu
Patrick Valduriez

Second Edition

分布式数据库 系统原理

(第 2 版)



清华大学出版社

Good luck!